

Burst-tolerance in Jellyfish

A New Routing Algorithm

By: Liwei Cui, Mou Zhang, Yifeng Yin



Content

1. Problem: burst traffic
2. Recap: Jellyfish
3. Jellyfish's Approach
4. Motivation
5. Our Progress
6. Next Steps
7. Acknowledgement


WHAT IS BURST TRAFFIC?

“Burst traffic”, is a sudden, uncertain, unpredictable traffic peak. Cause big trouble for network.




Recap: Jellyfish

Forget about structure – let's have no structure at all!



(a)




(b)



Recap: Jellyfish

Identical Equipment, **25%** more throughput!




“ While the above experiments establish that Jellyfish topologies have **high capacity**, it remains unclear whether this potential can be realized in real networks.

Section 5 in Jellyfish Paper

”



ECMP Does Not Work




1 → 6: only one(1) path is available; it uses two(2) links (1 → 5, 5 → 6)

WHAT IS JELLYFISH'S APPROACH?

“K-shortest path routing”. It is not only about a shortest path but also about next $k - 1$ shortest paths (longer, maybe).

3-shortest Paths Works (Much) Better




1 → 6: three(3) paths are available:

1. **1 → 5 → 6**
2. **1 → 7 → 8 → 6**
3. **1 → 7 → 9 → 6**



Path Diversity



For each **link**, we count the number of **distinct paths** it is on.

The more the path diversity, the less path overlap may occur.


Motivation

What if we maximize the path diversity, to achieve higher throughput?

WHAT IS OUR ALGORITHM?

“K-non-overlapping path routing”. It is guarantees all links on paths from A to B have no overlapping (longer, likely).

3-non-overlapping Paths Works Best




1 → 6: three(3) paths are available:

- 1. 1 → 5 → 6**
- 2. 1 → 7 → 8 → 6**
- 3. 1 → 2 → 3 → 4 → 6**

ECMP (Shortest Paths)

1 → 6: One(1) path is available:

1. **1 → 5 → 6**



3-shortest Paths

1 → 6: three(3) paths are available:

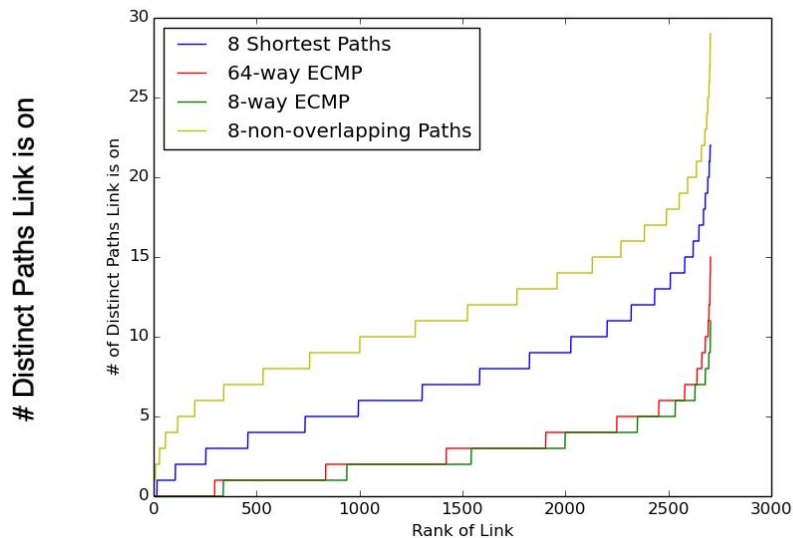
1. **1 → 5 → 6**
2. **1 → 7 → 8 → 6**
3. **1 → 7 → 9 → 6**

3-non-overlapping Paths

1 → 6: three(3) paths are available:

1. **1 → 5 → 6**
2. **1 → 7 → 8 → 6**
3. **1 → 2 → 3 → 4 → 6**


Path Diversity




For each **link**, we count the number of **distinct paths** it is on.

The more the path diversity, the less path overlap may occur.

switches layer



Host layer



Red is server
Gray is client



Experiment Configuration

50

switches

8

ports connecting peer
switches

10 Mbps

bandwidth between
switches

12

random hosts as
servers

12

random hosts as clients

100 Mbps

bandwidth between switches
and hosts

↑23%

8-Shortest-Paths achieving **27.98 Mbps** throughput
8-Non-overlapping achieving **34.58 Mbps** throughput



Our Progress

Reproduce Jellyfish

We leveraged several libraries (Mininet, Pox, RipL, RipL-POX) and open-source code to reproduce the Jellyfish network and k-shortest-paths routing

Comparison and Test

We compared the path diversity between non-overlapping routing with k-shortest-path & tested the average throughput per server, achieving 23.6% more throughput

New Routing Algorithm

We put forward and implemented a new routing algorithm (*Non-overlapping Path Algorithm*), which guarantees all links on paths from A to B have no overlapping

Deployment on Cloud

We set up our experimental environment on the Google Cloud Platform, which makes it easier to conduct future experiments.


“

**Talk is
cheap.
Show me
the code.**

Linus Torvalds

”

Non-overlapping-Path-in-Jellyfish @ Github

 Lw-Cui / Non-overlapping-Path-in-Jellyfish

Unwatch 1

Star 0

Fork 0

<> Code

Issues 0

Pull requests 0

Actions

Projects 0

Wiki

Security

Insights

Settings

Non-overlapping Path Algorithm for Jellyfish

Edit

jellyfish

network

routing-algorithm

k-shortest-paths

topology

Manage topics

15 commits

1 branch

0 packages

0 releases

1 contributor

Branch: master


New pull request






Create new file

Upload files

Find file

Clone or download

 Lw-Cui Update README.md Latest commit 8bc36fa 5 minutes ago

 pox	add all necessary file	4 days ago
 ripl	add all necessary file	4 days ago
 riplox	add all necessary file	4 days ago
 .gitignore	first init	6 days ago
 README.md	Update README.md	5 minutes ago

WHAT IS NEXT STEPS?



RELEASE RESTRAINT

Release strict non-overlapping restraint to balance the path length and path diversity;



MORE EXPERIMENT

Perform more tests under various topologies and circumstances.



Acknowledge

1. **Mininet** library for network emulation
2. **POX** library for OpenFlow controller
3. **RipL** library for simplifying data center code
4. **RipL-POX** library for controller built on RipL
5. **Austin Poore** and **Tommy Fan's** repo for inspiration to reproduce Jellyfish



Burst-tolerance in Jellyfish

Thank you for listening